

# Some R tricks

basically for more speed

Rochebrune – March 2016



The R inferno, Patrick Burns

<http://www.burns-stat.com/documents/books/the-r-inferno/>



FasterR! HigheR! StrongerR!, Noam Ross

<http://www.noamross.net/blog/2013/4/25/faster-talk.html>



Seamless R and C++ integration with Rcpp, Dirk Eddelbuettel

<http://dirk.eddelbuettel.com>



Hadley Wickham, ggplot2, an implementation of the grammar of graphics


<http://had.co.nz/>, <http://ggplot2.org/>, <http://yihui.name/knitr/>

# Part I

Benchmark your code

# How to quickly benchmark your code

```
func.one <- function(n) {return(rnorm(n,0,1))}  
func.two <- function(n) {return(rpois(n,1))}  
  
library(microbenchmark)  
n <- 1000  
res <- microbenchmark(func.one(n), func.two(n), times=1000)  
autoplot(res)
```



# How to profile your code I

Suppose you want to evaluate which part of the following function is hot:

```
## generate data, center/scale and perform ridge regression
my.func <- function(n,p) {

  require(MASS)

  ## draw data
  x <- matrix(rnorm(n*p),n,p)
  y <- rnorm(n)

  ## center/scale
  xs <- scale(x)
  ys <- y-mean(y)

  ## return ridge's coefficients
  ridge <- lm.ridge(ys~xs+0,lambda=1)

  return(ridge$coef)
}
```

# How to profile your code II

One can rely on the default Rprof function, with somewhat technical outputs

```
Rprof(file="profiling.out", interval=0.05)
res <- my.func(1000,500)
```

```
## Loading required package: MASS
```

```
Rprof(NULL)
```

```
summaryRprof("profiling.out")$by.self
```

##	self.time	self.pct	total.time	total.pct
## "La.svd"	1.00	74.07	1.00	74.07
## ".External2"	0.05	3.70	0.10	7.41
## "matrix"	0.05	3.70	0.10	7.41
## "aperm.default"	0.05	3.70	0.05	3.70
## "apply"	0.05	3.70	0.05	3.70
## "is.finite"	0.05	3.70	0.05	3.70
## "na.omit.data.frame"	0.05	3.70	0.05	3.70
## "rnorm"	0.05	3.70	0.05	3.70

# How to profile your code III


```
summaryRprof("profiling.out")$by.total
```

##	total.time	total.pct	self.time	self.pct
## "<Anonymous>"	1.35	100.00	0.00	0.00
## "block_exec"	1.35	100.00	0.00	0.00
## "call_block"	1.35	100.00	0.00	0.00
## "doTryCatch"	1.35	100.00	0.00	0.00
## "eval"	1.35	100.00	0.00	0.00
## "evaluate_call"	1.35	100.00	0.00	0.00
## "FUN"	1.35	100.00	0.00	0.00
## "handle"	1.35	100.00	0.00	0.00
## "in_dir"	1.35	100.00	0.00	0.00
## "knit"	1.35	100.00	0.00	0.00
## "lapply"	1.35	100.00	0.00	0.00
## "my.func"	1.35	100.00	0.00	0.00
## "process_file"	1.35	100.00	0.00	0.00
## "process_group"	1.35	100.00	0.00	0.00
## "process_group.block"	1.35	100.00	0.00	0.00
## "try"	1.35	100.00	0.00	0.00
## "tryCatch"	1.35	100.00	0.00	0.00
## "tryCatchList"	1.35	100.00	0.00	0.00
## "tryCatchOne"	1.35	100.00	0.00	0.00
## "withCallingHandlers"	1.35	100.00	0.00	0.00
## "withVisible"	1.35	100.00	0.00	0.00
## "lm.ridge"	1.15	85.19	0.00	0.00
## "svd"	1.05	77.78	0.00	0.00
## "La.svd"	1.00	74.07	1.00	74.07
## ".External2"	0.10	7.41	0.05	3.70
## "matrix"	0.10	7.41	0.05	3.70
## "scale"	0.10	7.41	0.00	0.00
## "scale.default"	0.10	7.41	0.00	0.00
## "aperm.default"	0.05	3.70	0.05	3.70
## "apply"	0.05	3.70	0.05	3.70
## "is.finite"	0.05	3.70	0.05	3.70

# How to profile your code III

The profr package is maybe a little easier to understand...

```
library(profr)
profiling <- profr({my.func(1000,500)}, interval=0.01)
plot(profiling)
```



## Part II

Use multiple cores for your simulation



# The `do.call` function

*constructs and executes a function call from a name or a function and a list of arguments to be passed to it*

Suppose you have the outputs of 100 simulations at your disposal, stored in a list like that

```
res[[1]]  
  
##      method      mse      timing  
## 1  lasso 0.7862968  0.9399695  
## 2  ridge 0.5057219  0.7958627  
## 3  bayes 0.9310022 115.8219670  
  
length(res)  
  
## [1] 100
```

How would you store them in a single data frame?

```
all.res <- do.call(rbind, res)  
dim(all.res)  
  
## [1] 300  3
```

# The `do.call` function

*constructs and executes a function call from a name or a function and a list of arguments to be passed to it*

Suppose you have the outputs of 100 simulations at your disposal, stored in a list like that

```
res[[1]]  
  
##      method      mse      timing  
## 1  lasso 0.7862968  0.9399695  
## 2  ridge 0.5057219  0.7958627  
## 3  bayes 0.9310022 115.8219670
```

```
length(res)
```

```
## [1] 100
```

How would you store them in a single data frame?

```
all.res <- do.call(rbind, res)  
dim(all.res)
```

```
## [1] 300  3
```

# Parallelizing is very easy I

Do some parallel computation as soon as you do simulations (this should happen sometimes)

```
library(parallel) ## embedded with R since version 2.9 or something  
cores <- detectCores() ## How many cores do I have?  
print(cores)
```

```
## [1] 4
```

My simulation study estimates the test error from ridge regression

```
one.simu <- function(i) {  
  ## draw data  
  n <- 1000; p <- 500  
  x <- matrix(rnorm(n*p),n,p) ; y <- rnorm(n)  
  ## return ridge's coefficients  
  train <- 1:floor(n/2)  
  test  <- setdiff(1:n,train)  
  ridge <- lm.ridge(y~x+0,lambda=1,subset=train)  
  err <- (y[test] - x[test, ] %*% ridge$coef )^2  
  return(list(err = mean(err), sd = sd(err)))  
}
```

# Parallelizing is very easy II


```
out <- mclapply(1:8, one.simu, mc.cores=cores)
head(do.call(rbind, out))
```

```
##      err      sd
## [1,] 13.72301 18.94939
## [2,] 10.66215 16.23548
## [3,]  9.22876 13.05196
## [4,]  9.438628 12.84848
## [5,] 10.57839 14.97044
## [6,] 12.95024 17.24075
```

## Be careful though...

- ▶ Parallelize piece of code complex enough
- ▶ Do not choose stupidly the number of cores

```
res <- microbenchmark(s1core = mclapply(1:8, one.simu, mc.cores=1),  
                      s2cores = mclapply(1:8, one.simu, mc.cores=2),  
                      s8cores = mclapply(1:8, one.simu, mc.cores=8), times=10)
```



# The Reduce function

*'Reduce' uses a binary function to successively combine the elements of a given vector*

↪ can be use to post-process your list of simulations obtained via `mclapply`

## Example


Work in progress with Avner for "jackknifing" a lasso solution path

```
rm(list=ls())  
library(lars)  
library(glmnet)  
## the diabetes data set (part of the lars package)  
data(diabetes)  
y <- diabetes$y  
x <- diabetes$x  
n <- length(y)
```

# The Reduce function II

## A single lasso fit

```
## recover a grid of lambda on the complete data set  
lasso <- glmnet(x,y)
```



# The Reduce function III

## Jackknifing the path

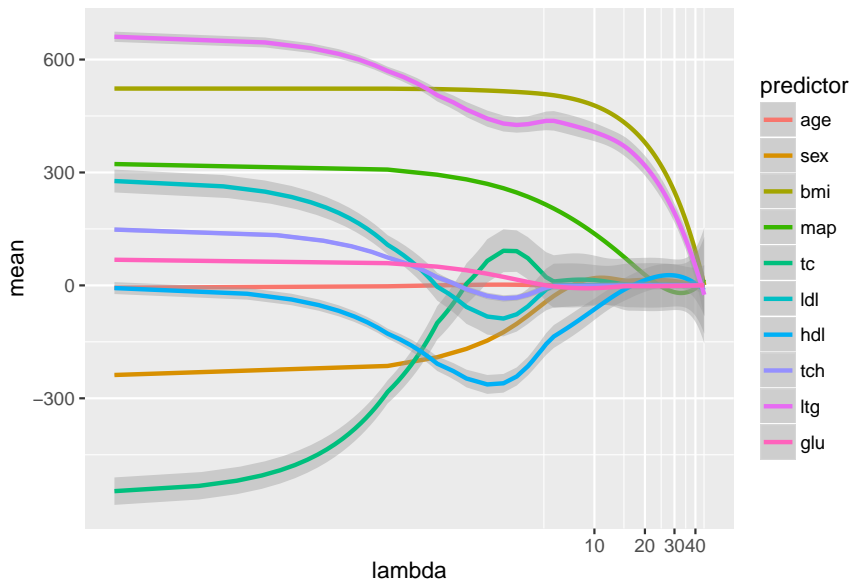
```
library(parallel)
## compute the regularization paths for all subsets,
## removing one individual at once
paths <- mclapply(1:n, function(i) {
  glmnet(x[-i, ], y[-i], lambda = lasso$lambda)$beta
}, mc.cores=4)
```

## Computing the envelop around the average regularization path with Reduce

```
mean.path <- Reduce("+", paths)/n
sdev.path <- sqrt(Reduce("+", lapply(paths, function(path) path**2))/n -
  mean.path**2)
```



# The Reduce function IV



## Part III

Be aware of what R is good (and bad) for

# Use the vector capabilities of R

Any algebraic operation should be thought in a “vectorized” way

```
exp2.1 <- sum(2^(0:10)/c(1,cumprod(1:10))) ## good
exp2.2 <- 1
for(k in 1:10) ## bad
  exp2.2 <- exp2.2 + 2^k/factorial(k)
```

Even non-algebraic operation should be thought as algebraic:

```
outer(1:4,c("A","B","C","D"),FUN=paste,sep="-")
```

```
##      [,1] [,2] [,3] [,4]
## [1,] "1-A" "1-B" "1-C" "1-D"
## [2,] "2-A" "2-B" "2-C" "2-D"
## [3,] "3-A" "3-B" "3-C" "3-D"
## [4,] "4-A" "4-B" "4-C" "4-D"
```

# Use the vector capabilities of R

Any algebraic operation should be thought in a “vectorized” way

```
exp2.1 <- sum(2^(0:10)/c(1,cumprod(1:10))) ## good
exp2.2 <- 1
for(k in 1:10) ## bad
  exp2.2 <- exp2.2 + 2^k/factorial(k)
```


Even non-algebraic operation should be thought as algebraic:

```
outer(1:4,c("A","B","C","D"),FUN=paste,sep="-")

##      [,1]  [,2]  [,3]  [,4]
## [1,] "1-A" "1-B" "1-C" "1-D"
## [2,] "2-A" "2-B" "2-C" "2-D"
## [3,] "3-A" "3-B" "3-C" "3-D"
## [4,] "4-A" "4-B" "4-C" "4-D"
```

# Preallocate whenever it is possible

```
grow <- function(n) {vec <- numeric(0); for (i in 1:n) vec <- c(vec,i)}  
loop <- function(n) {vec <- numeric(n); for (i in 1:n) vec[i] <- i}  
vect <- function(n) {1:n}
```



# Do not stack objects I

Even if it is tempting when the final size is unknown.


```
simu.stack <- function(x) { ## x is a n x p matrix
  out <- data.frame(mean = numeric(0), sd = numeric(0))
  for (i in 1:n)
    out <- rbind(out, data.frame(mean = mean(x[i,]), sd = sd(x[i, ])) )
  return(out)
}

simu.df <- function(x) {
  out <- data.frame(mean = numeric(n), sd = numeric(n))
  for (i in 1:n)
    out[i, ] <- c(mean = mean(x[i,]), sd = sd(x[i, ]))
  return(out)
}

simu.list <- function(x) {
  my.list <- lapply(1:n, function(i) c(mean(x[i,]), sd(x[i, ])))
  out <- data.frame(do.call(rbind, my.list))
  colnames(out) <- c("mean", "sd")
  return(out)
}
```

# Do not stack objects II

```
n <- 1000; p <- 10; x <- matrix(rnorm(n*p), n, p)
res <- microbenchmark(simu.stack(x), simu.df(x), simu.list(x), times=20)
```



# Use the [a-z]\*pply family

## Example with factors (tapply)

```
data <- rnorm(100)
sexe <- factor(sample(c("H", "F"), 100, rep=TRUE))
mean.1 <- tapply(data, sexe, mean) ## good
mean.2 <- c() ## complicated
for (l in levels(sexe))
  mean.2 <- c(mean.2, mean(data[sexe == l]))
```

## Example with list or data.frame (sapply/lapply)

```
data(oats)
oats[1:2, ]
```

```
##      B      V      N      Y
## 1 I Victory 0.0cwt 111
## 2 I Victory 0.2cwt 130
```

```
sapply(oats, is.factor) ## readable
```

```
##      B      V      N      Y
## TRUE TRUE TRUE FALSE
```

```
for (c in 1:ncol(oats)) ## less readable (I think)
  print(is.factor(oats[,c]))
```



# Use the [a-z]\*pply family

## Example with factors (tapply)

```
data <- rnorm(100)
sexe <- factor(sample(c("H", "F"), 100, rep=TRUE))
mean.1 <- tapply(data, sexe, mean) ## good
mean.2 <- c() ## complicated
for (l in levels(sexe))
  mean.2 <- c(mean.2, mean(data[sexe == l]))
```

## Example with list or data.frame (sapply/lapply)

```
data(oats)
oats[1:2, ]

##      B      V      N      Y
## 1 I Victory 0.0cwt 111
## 2 I Victory 0.2cwt 130

sapply(oats, is.factor) ## readable


##      B      V      N      Y
## TRUE TRUE TRUE FALSE

for (c in 1:ncol(oats)) ## less readable (I think)
  print(is.factor(oats[,c]))
```

# Compile your functions I

```
cumsum.R <- function(n) {  
  x <- rnorm(n)  
  return(cumsum(x))  
}  
  
cumsum.me <- function(n) {  
  x <- rnorm(n)  
  res <- 0  
  for (i in 1:length(x)) {  
    res <- res+x[i]  
  }  
  return(res)  
}  
  
library(compiler)  
cumsum.cmp <- cmpfun(cumsum.me)  
  
n <- 1000  
res <- microbenchmark(cumsum.R(n), cumsum.me(n), cumsum.cmp(n), times=1000)
```


# Compile your functions II



# The crossprod function

As can be guessed, it computes the cross-product between two vector or matrices... and is generally fastest than `% * %` !


```
crossprod.prod <- function(x) return(t(x) %*% x)
crossprod.func1 <- function(x) return(crossprod(x,x))
crossprod.func2 <- function(x) return(crossprod(x))
```



# The row/colSums family

col/rowSums, col/rowMeans and their extensions in the matrixStats package (rank,max,min, etc.) are very efficient.

```
colSums.default <- function(x) return(colSums)
colSums.algebra <- function(x) return(crossprod(rep(1,nrow(x)), x))
colSums.apply <- function(x) return(apply(x,2,sum))
colSums.loop <- function(x) {
  res <- rep(0,ncol(x))
  for (i in 1:ncol(x))
    res[i] <- sum(x[,i])
  return(res)
}
```




# The secret function rowsum

`rowsum` (not to be confused with `rowSums`) computes sums in a vector split according a grouping variable (work for matrices).

```
vec <- runif(1000)
grp <- sample(1:5, 1000, TRUE)
print(c(rowsum(vec, grp)))
```

```
## [1] 102.87880  99.03421  93.72513  97.05151  89.06845
```


```
res <- microbenchmark(
  rowsum = rowsum(vec, grp),
  split  = sapply(split(vec, grp), sum),
  tapply = tapply(vec, grp, sum),
  aggreg = aggregate(vec, list(grp), sum),
  times=1000)
```



# Internal function are faster

Function defined internally are sometimes incredibly faster (written in C), but cannot be called in packages submitted to CRAN.

```
x <- rnorm(100)
res <- microbenchmark(mean = mean(x),
  mine = sum(x)/length(x),
  internal = .Internal(mean(x)), times = 1e+05)
```




# The sweep function

Fancy way to apply a statistic on a given dimension of an array.

```
center1 <- function(x) return(scale(x, colMeans(x), FALSE))
center2 <- function(x) return(sweep(x, 2, colMeans(x), "-", check.margin=FALSE))
center3 <- function(x) return(x - outer(rep(1, nrow(x)), colMeans(x)) )

seq.p <- 10^(1:5); n <- 100; times <- 20
```






# Basic algebra does not always pay

## Example for scaling a matrix

```
scale1 <- function(x) return(scale(x, FALSE, colSums(x^2)))  
scale2 <- function(x) return(sweep(x, 2, colSums(x^2), "/", check.margin=FALSE))  
scale3 <- function(x) return(x %*% diag(1/colSums(x^2)) )  
  
seq.p <- 10^(1:3); n <- 100; times <- 20
```



# Mind some algebra I


Example by inverting a positive definite matrices

```
use.chol <- function(n,p) {  
  x <- matrix(rnorm(n*p),n,p)  
  xtx <- crossprod(x)  
  return(chol2inv(chol(xtx)))  
}  
  
use.solve <- function(n,p) {  
  x <- matrix(rnorm(n*p),n,p)  
  xtx <- crossprod(x)  
  return(solve(xtx))  
}  
  
bench.p.fixed <- function(p, times) {  
  res <- microbenchmark(solve = use.solve(2*p,p),  
                        chol   = use.chol (2*p,p), times=times)  
  return(data.frame(method = res$expr,  
                    timings = res$time,  
                    size    = rep(as.character(p),times)))  
}
```

# Mind some algebra II

```
out <- do.call(rbind,  
              lapply(c(10,50,100,250),  
                    bench.p.fixed, times=10)  
              )  
  
head(out)  
  
##    method timings size  
## 1    chol  239439    10  
## 2   solve  786924    10  
## 3   solve   68059    10  
## 4   solve   52322    10  
## 5   solve   49488    10  
## 6    chol   45483    10  
  
p <- ggplot(out, aes(x=size, y=timings, fill=method)) +  
  geom_boxplot() + coord_trans(y="log10")
```

# Mind some algebra III



## Part IV

Remind that R is object oriented

## R masks the numerical errors

by printing a convenient summary of objects

```
7/13
```

```
## [1] 0.5384615
```

```
print(7/13, digits=16)
```

```
## [1] 0.5384615384615384
```

So do not use binary operator to compare floats because

```
.1 == (.3/3)
```

```
## [1] FALSE
```

```
print(.3/3, digits=16)
```

```
## [1] 0.09999999999999999
```

Try

```
all.equal(.1, .3/3)
```

```
## [1] TRUE
```

## R masks the numerical errors

by printing a convenient summary of objects

```
7/13  
  
## [1] 0.5384615  
  
print(7/13, digits=16)  
  
## [1] 0.5384615384615384
```

So do not use binary operator to compare floats because

```
.1 == (.3/3)  
  
## [1] FALSE  
  
print(.3/3, digits=16)  
  
## [1] 0.09999999999999999
```

Try

```
all.equal(.1, .3/3)  
  
## [1] TRUE
```

## R masks the numerical errors

by printing a convenient summary of objects

```
7/13
```

```
## [1] 0.5384615
```

```
print(7/13, digits=16)
```

```
## [1] 0.5384615384615384
```

So do not use binary operator to compare floats because

```
.1 == (.3/3)
```

```
## [1] FALSE
```

```
print(.3/3, digits=16)
```

```
## [1] 0.09999999999999999
```

Try

```
all.equal(.1, .3/3)
```

```
## [1] TRUE
```




# Factor conversion are slow (nlevels)

Do not use factor if you need to perform just one operation on it.

```
nlevels.factor <- function(n,K) {  
  x <- sample(1:K, n, rep=TRUE)  
  return(nlevels(factor(x)))  
}
```

```
nlevels.numeric <- function(n,K) {  
  x <- sample(1:K, n, rep=TRUE)  
  return(length(unique(x)))  
}
```

```
res <- microbenchmark(factor = nlevels.factor (1000,10),  
                      numeric = nlevels.numeric(1000,10), times=1000)
```




# Operations on factors are fast (nlevels)

Use factor if you need repeated operations on the same vector.

```
nk <- 20
seq.K <- c(10,100,1000)
res <- do.call(rbind, lapply(seq.K, function(K) {
  x1 <- rep(1:K,nk)
  x2 <- factor(x1)
  out <- microbenchmark(factor = nlevels(x2),
                        numeric = length(unique(x1)), times=1000)
  return(data.frame(method = out$expr, timings = out$time, K = factor(K)))
})))
```

```
head(res)
```


##	method	timings	K
## 1	factor	15505	10
## 2	numeric	20432	10
## 3	numeric	7436	10
## 4	numeric	6203	10
## 5	factor	3353	100
## 6	factor	2521	100



# Avoid table whenever you can

table is a complex function that should not be use for simple operations like counting the occurrences of integers in a vector.

```
n <- 1000  
K <- 10  
res <- microbenchmark(table      = table      (sample(1:K, n, rep=TRUE)),  
                        tabulate   = tabulate   (sample(1:K, n, rep=TRUE)),  
                        times=1000)
```



# Use the Matrix package I


Propose a collection of functions for of matrix algebra adapted to the type of matrix at hand (sparse, diagonal, triangular, block diagonal, etc.)

```
library(Matrix)
bench.par.fixed <- function(par) {
  n <- par$n; density <- par$density
  data <- sample(c(0,1),n**2,rep=TRUE,prob=c(1-density,density))
  x.dense <- matrix(data,n,n)
  x.sparse <- Matrix(data,n,n)
  res <- microbenchmark(dense = crossprod(x.dense) ,
                        sparse = crossprod(x.sparse), times=10)
  return(data.frame(method = res$expr,
                    timings = res$time,
                    size = n ,
                    density = density ))
}
```

# Use the Matrix package II

```
par <- expand.grid(n=c(100,250,500,1000), density=c(.05,.2,.5,.9))
out <- do.call(rbind,
  lapply(1:nrow(par),
    function(k) {
      return(bench.par.fixed(par[k, ]))
    })
)
out$density <- factor(out$density)
out$method <- factor(out$method)
out$size <- factor(out$size)
p <- ggplot(out, aes(x=size, y=timings, fill=method)) +
  geom_boxplot() + coord_trans(y="log10") + facet_wrap(~density, nrow=2)
```

# Use the Matrix package III



## Part V

Use (supposedly) lower-level languages

# Interfacing C++ with R is really easy I

## Example 1

For a vector  $\mathbf{x} = (x_1, \dots, x_n)$ , consider the simple task of computing

$$y_k = \sum_{i=1}^k \log(x_i), \quad k = 1, \dots, n.$$

One can easily integrate some C++ version of this code with Rcpp.

```
library(Rcpp)
cppFunction('NumericVector rcpp(NumericVector x) {
    using namespace Rcpp;

    int n = x.size() ;
    NumericVector res(x) ;
    res(0) = log(x(0));
    for (int i=1; i<n; i++) {
        res(i) = res(i-1) + log(x(i)) ;
    }
    return(wrap(res)) ;
}')
```



# Interfacing C++ with R is really easy II

## Example 1


The Armadillo library for linear algebra facilitates even more the integration

```
cppFunction(depends="RcppArmadillo", 'NumericVector Arma(NumericVector x) {  
  using namespace Rcpp;  
  using namespace arma;  
  return(wrap(cumsum(log(as<vec>(x))))) ;  
'})
```

# Interfacing C++ with R is really easy III

## Example 1

```
x <- runif(1e7, 1,2)
res <- microbenchmark(cpp = rcpp(x), arma = Arma(x), times=40)
```



# Interfacing C++ with R is really easy

Example 2: from a work with C. Lévy-Leduc and V. Brault

Let  $\mathbf{T}$  be an  $n \times n$  lower triangular matrix with nonzero elements equal to one. We need fast computation of

$$\text{vec}(\mathbf{TBT}^\top) = (\mathbf{T} \otimes \mathbf{T}) \times \text{vec}(\mathbf{B}).$$

```
library(Matrix); library(inline); library(RcppArmadillo)

prod.rough <- function(B) {
  n <- ncol(B); T <- bandSparse(n,k=(-n+1):0)
  return(kronecker(T,T) %*% as.vector(B))}

prod.smart <- function(B) {
  return(as.vector(apply(apply(B,1,cumsum),1,cumsum)))}

prod.wise <- cxxfunction(signature(B="matrix"), '
  using namespace Rcpp;
  using namespace arma;
  return(wrap(vectorise(cumsum(cumsum(as<mat>(B),0),1)))) ;
  ', plugin="RcppArmadillo")
```

# Interfacing C++ with R is really easy II

## Example 2

```
B.ls <-  apply(c(10, 50, 100), function(n) matrix(runif(n**2),n,n))
res <- do.call(rbind, lapply(B.ls, function(B) {
  out <- microbenchmark(rough = prod.rough(B),
                        smart = prod.smart(B),
                        wise  = prod.wise(B), times=20)
  return(data.frame(method = out$expr, timings = out$time, n = factor(ncol(B))))
})))
```

